

# **Lab notebook Week 9**

**Submitted by: Vishrut Sharma (OdinID: vishrut)**

## **Table of Contents**

<b>09.1g: BigQuery, BigLake.....</b>	<b>1</b>
3. Create dataset.....	1
4. Query data.....	2
9. Query data.....	3
<b>09.2g: Jupyter Notebooks.....</b>	<b>3</b>
3. BigQuery query.....	3
6. Run queries.....	4
8. Mobility.....	4
9. Airport traffic.....	4
10. Mortality.....	5
11. Run example queries.....	5
12. Write queries.....	7
<b>09.3g: Dataproc.....</b>	<b>8</b>
6. Run computation.....	8
8. Run computation again.....	8
<b>09.4g: Dataflow.....</b>	<b>8</b>
3. Beam code.....	8
4. Run pipeline locally.....	9
5. Dataflow Lab #2 (Word count).....	10
6. Run code locally.....	11
9. Run code using Dataflow runner.....	11
12. View raw data from PubSub.....	12
14. Run Dataflow job from template.....	13
15. Query data in BigQuery.....	13
16. Data visualization.....	15

## 09.1g: BigQuery, BigLake

### 3. Create dataset

- Take a screenshot of the table's details that includes the number of rows in the table.

The screenshot shows the Google Cloud BigQuery console interface. The left sidebar displays the Explorer view with a search bar and a list of workspace resources. The main panel shows the details for the table 'yob\_native\_table'. The 'DETAILS' tab is active, displaying the following information:

- Schema:** Case insensitive: false
- Description:**
- Labels:**
- Primary key(s):**
- Storage info:**
  - Number of rows: 33,044
  - Total logical bytes: 618.78 KB
  - Active logical bytes: 618.78 KB
  - Long term logical bytes: 0 B
  - Total physical bytes: 0 B
  - Active physical bytes: 0 B
  - Long term physical bytes: 0 B
  - Time travel physical bytes: 0 B

A notification at the bottom states: "yob\_native\_table" created. Below the notification, there are buttons for "GO TO TABLE" and "Open Editor".

### 4. Query data

- Screenshot the query results and include it in your lab notebook

The screenshot shows the Google Cloud BigQuery console interface with a query executed. The query is displayed in the editor:

```
1 SELECT name, count FROM "cloud-sharma-vishrut.yob_native_table" where gender='F' order by count desc limit 20;
```

The query results are displayed in a table with the following data:

Row	name	count
1	Emma	20799
2	Olivia	19674
3	Sophia	18490
4	Isabella	16950
5	Ava	15586
6	Mia	13442
7	Emily	12562
8	Abigail	11985
9	Madison	10247
10	Charlotte	10048
11	Harper	9564
12	Sofia	9542
13	Avery	9517

The interface also shows tabs for "JOB INFORMATION", "RESULTS", "JSON", "EXECUTION DETAILS", and "EXECUTION GRAPH". The "RESULTS" tab is currently selected. At the bottom, there are buttons for "SAVE RESULTS", "EXPLORE DATA", and "REFRESH".

- Screenshot your results and include it in your lab notebook

```
vishrut@cloudshell:~ (cloud-sharma-vishrut)$ bq query "SELECT name, count from [cloud-sharma-vishrut.yob.yob_native_table] WHERE gender='M' ORDER BY count ASC LIMIT 10"
10530 22:36:26.733797 140627648767808 bigquery_client.py:730] There is no apilog flag so non-critical logging is disabled.
+-----+-----+
| name | count |
+-----+-----+
| Aari | 5 |
| Aaliyah | 5 |
| Aadian | 5 |
| Aaroh | 5 |
| Aarit | 5 |
| Aadiv | 5 |
| Aadhi | 5 |
| Aarohan | 5 |
| Aariyan | 5 |
| Aamer | 5 |
+-----+-----+
vishrut@cloudshell:~ (cloud-sharma-vishrut)$
```

- Screenshot your results and include it in your lab notebook

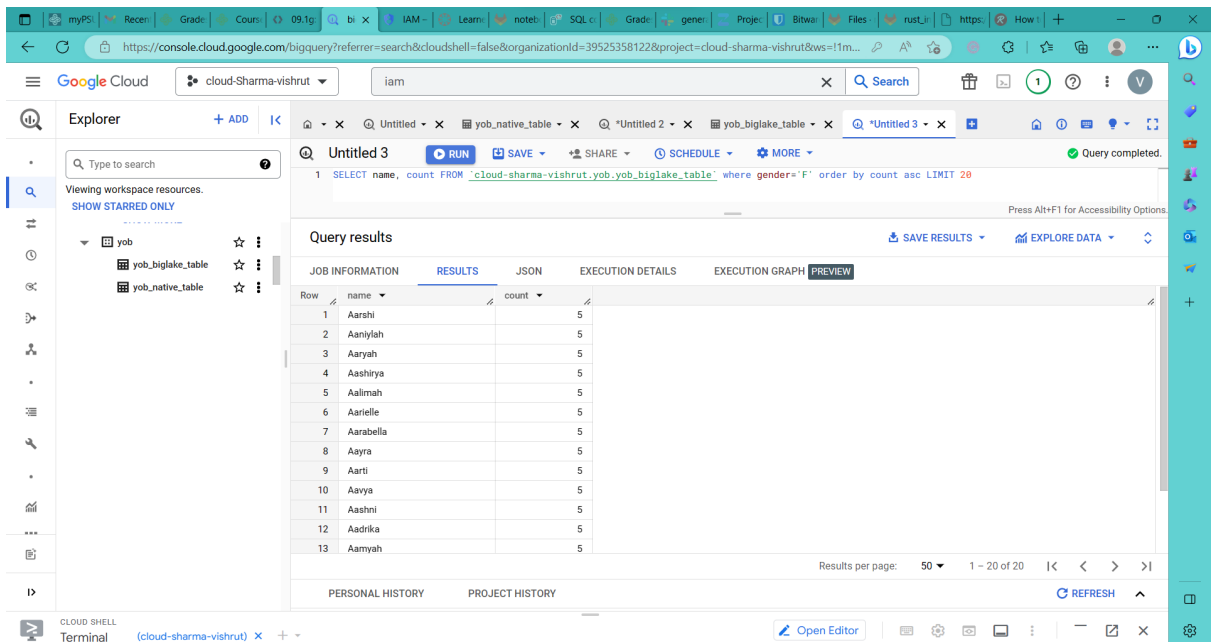
```
vishrut@cloudshell:~ (cloud-sharma-vishrut)$ bq shell
10530 22:37:51.507191 140689267836736 bigquery_client.py:730] There is no apilog flag so non-critical logging is disabled.
Welcome to BigQuery! (Type help for more information.)
cloud-sharma-vishrut> select name, count from [cloud-sharma-vishrut.yob.yob_native_table] where gender='M' order by count asc limit 10
+-----+-----+
| name | count |
+-----+-----+
| Aari | 5 |
| Aaliyah | 5 |
| Aadian | 5 |
| Aaroh | 5 |
| Aarit | 5 |
| Aadiv | 5 |
| Aadhi | 5 |
| Aarohan | 5 |
| Aariyan | 5 |
| Aamer | 5 |
+-----+-----+
cloud-sharma-vishrut>
```

- Screenshot your results and include it in your lab notebook

```
cloud-sharma-vishrut> select name, count from [cloud-sharma-vishrut.yob.yob_native_table] where name='Vishrut'
cloud-sharma-vishrut>
```

## 9. Query data

- Screenshot the query results and include it in your lab notebook



The screenshot shows the Google Cloud BigQuery console interface. The query editor displays the following SQL query:

```
1 SELECT name, count FROM 'cloud-sharma-vishrut.yob.yob_biglake_table' where gender='F' order by count asc LIMIT 20
```

The query results are displayed in a table with 13 rows and 2 columns: name and count. The results are ordered by count in ascending order.

Row	name	count
1	Aarshi	5
2	Aaniylah	5
3	Aaryah	5
4	Aashirya	5
5	Aalimah	5
6	Aarielle	5
7	Aarabella	5
8	Aayra	5
9	Aarti	5
10	Aavya	5
11	Aashni	5
12	Adrika	5
13	Aamyah	5

The interface also shows the Explorer panel on the left with the project structure, and the bottom panel with the Cloud Shell terminal.

## 09.2g: Jupyter Notebooks

### 3. BigQuery query

- How much less data does this query process compared to the size of the table?

Answer: 18.89 GB less data

- How many twins were born during this time range?

Answer: 375362

- How much lighter on average are they compared to single babies?

Answer: 2.17116045789 lighter



### 6. Run queries

- Show the plots generated for the two most important features for your lab notebook

Answer:


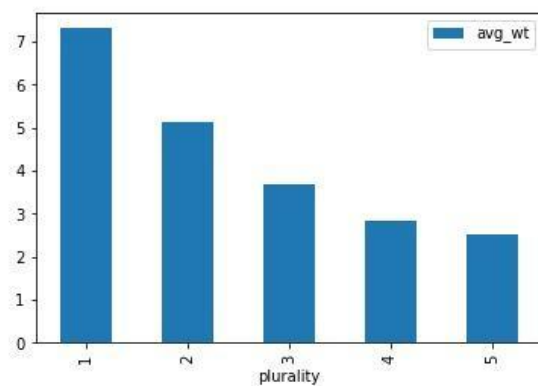
```
[8]: df = get_distinct_values('gestation_weeks')
df.plot(x='gestation_weeks', y='avg_wt', kind='bar')
```

[8]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f93ea6d5b50>



```
[7]: df = get_distinct_values('plurality')
df.plot(x='plurality', y='avg_wt', kind='bar')
```

[7]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f9448085b10>



## **8. Mobility**

- **What day saw the largest spike in trips to grocery and pharmacy stores?**

**Answer:** 2020-03-13

- **On the day the stay-at-home order took effect (3/23/2020), what was the total impact on workplace trips?**

**Answer:** -49

## **9. Airport traffic**

- **Which three airports were impacted the most in April 2020 (the month when lockdowns became widespread)?**

**Answer:** Newark Liberty International, Daniel K. Inouye International, Chicago OHare International

- **Run the query again using the month of August 2020. Which three airports were impacted the most?**

**Answer:** Newark Liberty International, Charlotte Douglas International, Dallas/Fort Worth International

## **10. Mortality**

- **What table and columns identify the place name, the starting date, and the number of excess deaths from COVID-19?**

**Answer:** Table name: excess\_deaths

Columns: placename, start\_date, excess\_deaths

- **What table and columns identify the date, county, and deaths from COVID-19?**

**Answer:** Table name: us\_counties

Columns: date, county, deaths

- **What table and columns identify the date, state, and confirmed cases of COVID-19?**

**Answer:** Table name: us\_states

Columns: date, state\_name, confirmed\_cases

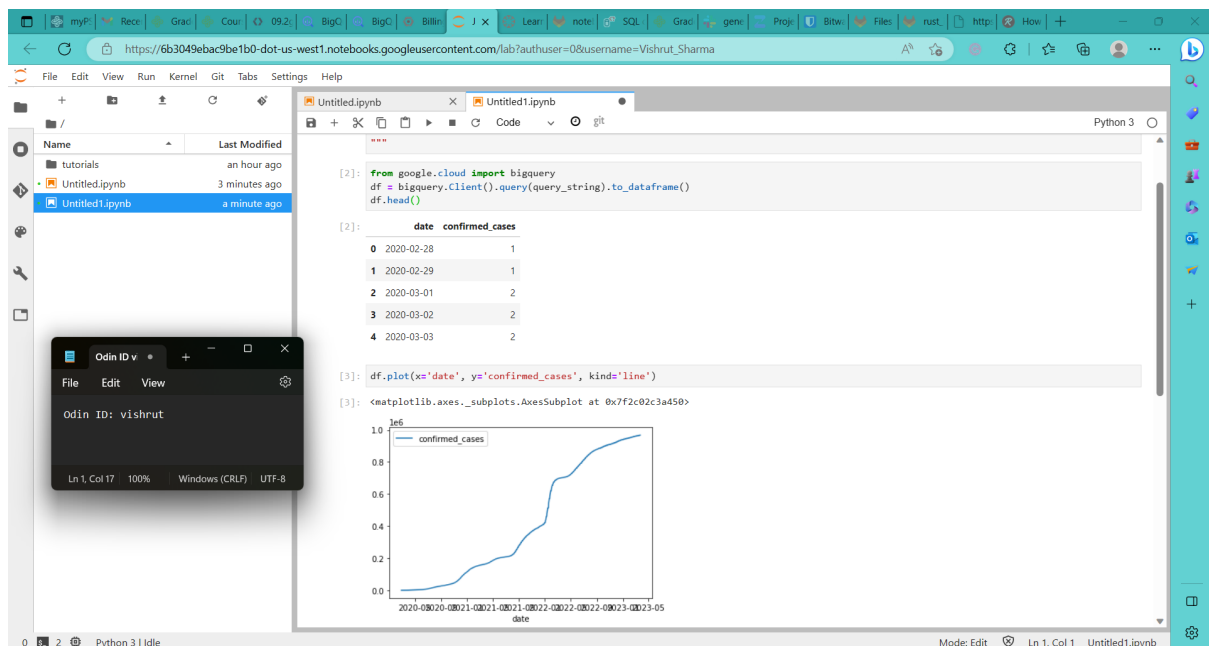
- **What table and columns identify a county code and the percentage of its residents that report they always wear masks?**

**Answer:** Table name: mask\_use\_by\_county

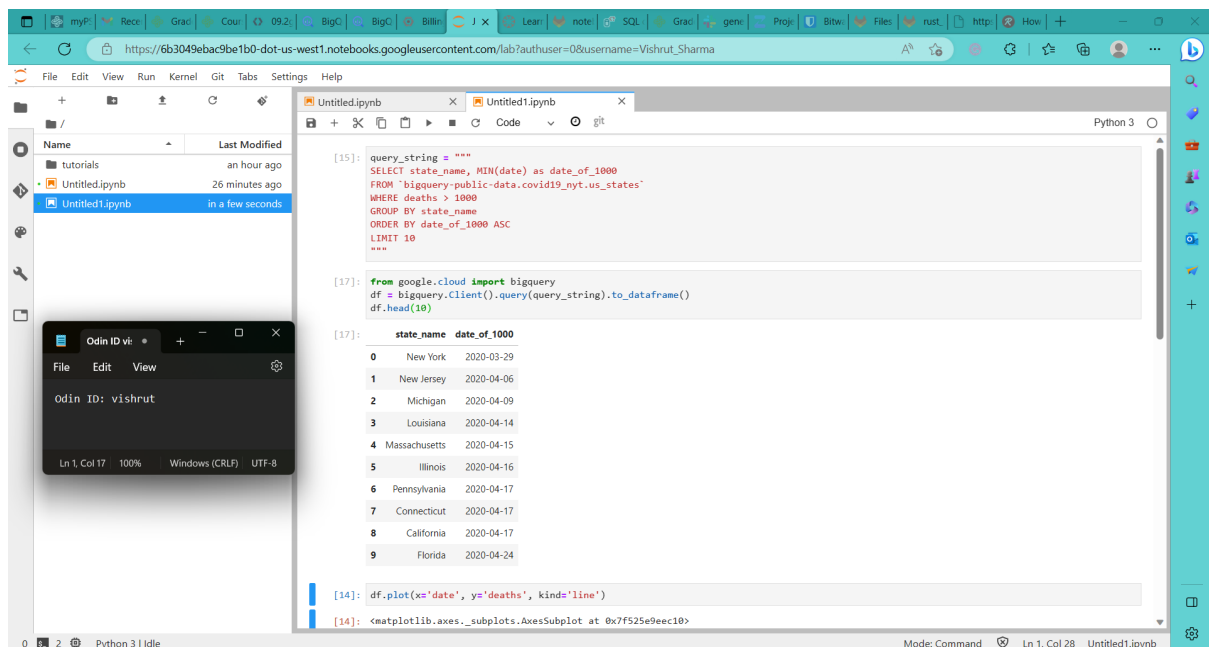
Columns: county\_fips\_code, always

## 11. Run example queries

- Show a screenshot of the plot and the code used to generate it for your lab notebook



- From within your Jupyter notebook, run the query and write code that shows the first 10 states that reached 1000 deaths from COVID-19. Take a screenshot for your lab notebook.



- Take a screenshot for your lab notebook of the Top 5 counties and the states they are located in.

The screenshot shows a Jupyter Notebook with two tabs: 'Untitled.ipynb' and 'Untitled1.ipynb'. The 'Untitled1.ipynb' tab is active, displaying a SQL query and its results.

**Code Cell [5]:**

```
query_string = """
SELECT DISTINCT mu.county_fips_code, mu.always, ct.county
FROM `bigquery-public-data.covid19_nyt.mask_use_by_county` as mu
LEFT JOIN `bigquery-public-data.covid19_nyt.us_counties` as ct
ON mu.county_fips_code = ct.county_fips_code
ORDER BY mu.always DESC
"""
```

**Code Cell [6]:**

```
from google.cloud import bigquery
df = bigquery.Client().query(query_string).to_dataframe()
df.head()
```

**Output of Cell [6]:**

	county_fips_code	always	county
0	06027	0.889	Inyo
1	36123	0.884	Yates
2	48229	0.880	Hudspeth
3	06051	0.880	Mono
4	48141	0.877	El Paso

## 12. Write queries

- Plot the results and take a screenshot for your lab notebook.

The screenshot shows a Jupyter Notebook with two tabs: 'Untitled.ipynb' and 'Untitled1.ipynb'. The 'Untitled1.ipynb' tab is active, displaying a SQL query, its results, and a line plot.

**Code Cell [7]:**

```
query_string = """
SELECT date, deaths
FROM `bigquery-public-data.covid19_nyt.us_counties`
WHERE county = 'Multnomah' AND state_name = 'Oregon'
ORDER BY date ASC
"""
```

**Code Cell [10]:**

```
from google.cloud import bigquery
df = bigquery.Client().query(query_string).to_dataframe()
```

**Code Cell [11]:**

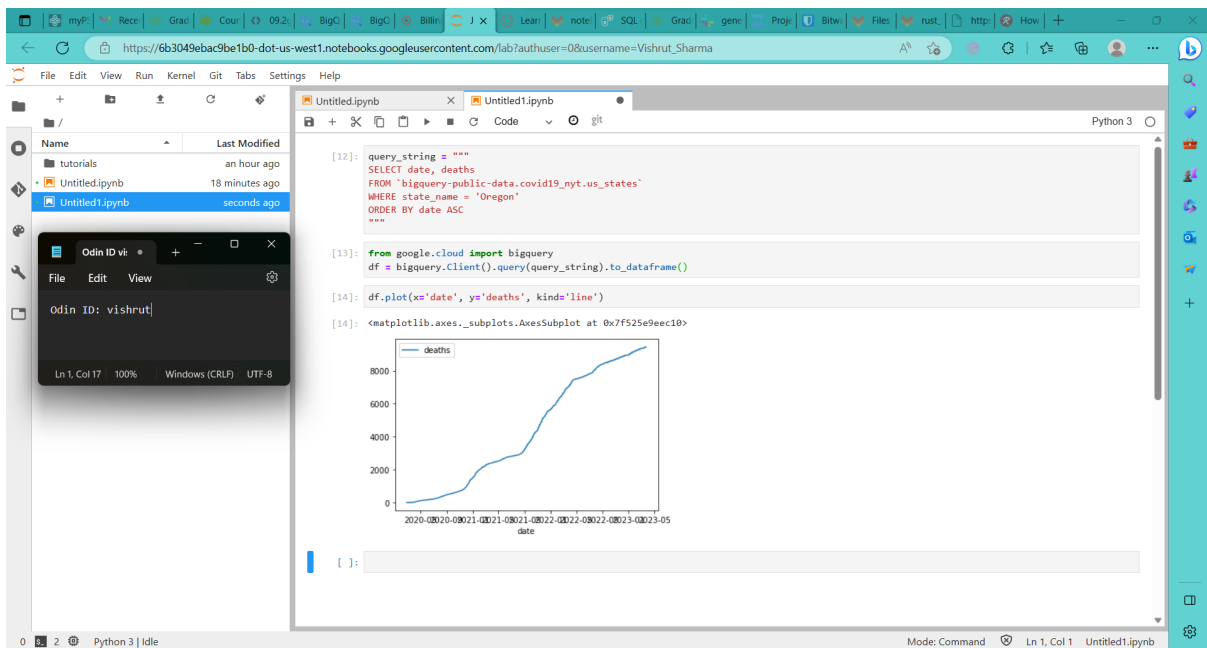
```
df.plot(x='date', y='deaths', kind='line')
```

**Output of Cell [11]:**

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f525f641d90>

The plot shows the number of deaths over time for Multnomah County, Oregon. The x-axis is labeled 'date' and ranges from 2020-09-01 to 2023-01-01. The y-axis is labeled 'deaths' and ranges from 0 to 1400. The plot shows a steady increase in deaths over time, with a significant jump around 2021-09-01.

- Plot the results and take a screenshot for your lab notebook.



## 09.3g: Dataproc

### 6. Run computation

- How long did the job take to execute?

Answer: Approximately 30 seconds

- Examine output.txt and show the estimate of  $\pi$  calculated.

Answer: 3.141495591414956

### 8. Run computation again

- How long did the job take to execute? How much faster did it take?

Answer: Approximately 19 seconds. It was 11 seconds faster.

- Examine output2.txt and show the estimate of  $\pi$  calculated.

Answer: 3.1415917714145079

## 09.4g: Dataflow

### 3. Beam code

- Where is the input taken from by default?

Answer: input is taken from

'./javahelp/src/main/java/com/google/cloud/training/dataanalyst/javahelp/'



- Where does the output go by default?

**Answer:** '/tmp/output'

- Examine both the `getPackages()` function and the `splitPackageName()` function. What operation does the 'PackageUse()' transform implement?

**Answer:** The `PackageUse()` transform takes a line of code and a keyword and returns a tuple with a count of 1 for each package used.

- Look up Beam's `CombinePerKey`. What operation does the `TotalUse` operation implement?

**Answer:** The `TotalUse` operation implements the Reduce operation using the `CombinePerKey` transform. It combines the counts for each packet and summarizes them based on the packet key.

- Which operations correspond to a "Map"?

**Answer:** `GetImports`, `PackageUse`

- Which operation corresponds to a "Shuffle-Reduce"?

**Answer:** `TotalUse`

- Which operation corresponds to a "Reduce"?

**Answer:** `Top_5`

## 4. Run pipeline locally

- Take a screenshot of its contents

```
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$ python is_popular.py
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$ ls
env  grepc.py  grep.py  install_packages_OLD.sh  install_packages.sh  is_popular.py  JavaProjectsThatNeedHelp.py  JavaProjectsThatNeedHelp_PY2_Version.py  OLD_grepc.py
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$ cd /tmp
-bash: cd /tmp: No such file or directory
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$ cd /tmp
(env) vishrut@cloudshell:/tmp (cloud-sharma-vishrut)$ ls
minikube_delete_53795bc2289bd7e8fbb94c747e1b59ccb36db06a_0.log  output-000000-of-00001  tmp.9WpEupIrfq  tmp.ETYuZnOU1H  tmpux-1000
(env) vishrut@cloudshell:/tmp (cloud-sharma-vishrut)$ cat output-000000-of-00001
[('org', 45), ('org.apache', 44), ('org.apache.beam', 44), ('org.apache.beam.sdk', 43), ('org.apache.beam.sdk.transforms', 16)]
(env) vishrut@cloudshell:/tmp (cloud-sharma-vishrut)$
```

- Explain what the data in this output file corresponds to based on your understanding of the program.

**Answer:** The output file will contain the top 5 Java packages that are most frequently used in the input Java code files. Each line in the output file will consist of a package name followed by its corresponding count, separated by a tab or space.

The program scans through the input Java code files and identifies lines starting with the keyword 'import'. It extracts the package names from those lines and keeps track of the occurrences of each package. The counts are combined using the 'CombinePerKey' transform.

The 'Top\_5' transform selects the 5 packages with the highest counts, based on their occurrence values. Finally, the selected packages and their counts are written to the output file.

Consequently, the output file provides valuable insights into the commonly used Java packages in the analyzed codebase. This information helps developers identify frequently employed dependencies and potentially optimize their code or manage dependencies more effectively.

## **5. Dataflow Lab #2 (Word count)**

- **What are the names of the stages in the pipeline?**

**Answer:** Read, Split, PairWithOne, GroupAndSum, Format, Write,

- **Describe what each stage does.**

**Answer:**

### **1. Read:**

- Utilizes the ReadFromText function to read the specified text file mentioned in the input argument.
- Generates a PCollection named "lines" that contains the individual lines from the input text file.

### **2. Split:**

- Applies the WordExtractingDoFn function using ParDo.
- Divides each line of text into separate words using regular expression matching.
- Produces a PCollection consisting of the individual words.

### **3. PairWithOne:**

- Maps each word to a key-value pair where the word becomes the key and the value is set as 1.
- Assigns an initial count of 1 to each word.

### **4. GroupAndSum:**

- Groups the key-value pairs by their keys (words).
- Combines the values (counts) associated with each key (word) using the CombinePerKey transform along with the sum function.
- Generates a PCollection of key-value pairs where each key represents a unique word, and the value is the sum of counts for that word.

### **5. Format:**

- Maps each key-value pair to a formatted string using the format\_result function.
- Formats the word and its count using the pattern "%s: %d", where %s represents the word and %d represents the count.

- Produces a PCollection of formatted strings.

## 6. Write:

- Uses the WriteToText transform to write the output PCollection to the output file specified in the --output argument.
- Each stage in the pipeline carries out a specific transformation or action on the data, such as reading from a text file, splitting lines into words, counting word occurrences, formatting results, and writing the output to a file.

## 6. Run code locally

- Use wc with an appropriate flag to determine the number of unique words in King Lear.

```
0x7f9178a67d60> for environment ref Environment default_environment_1 (beam:env:embedded_python:vi, b'')
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:apache_beam.io.filebasedsink:Starting finalize write threads with num_shards: 1 (skipped: 0), batches: 1, num_threads: 1
INFO:apache_beam.io.filebasedsink:Renamed 1 shards in 0.01 seconds.
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$ cat outputs-*--of-* | wc -w
9568
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$
```

- Use sort with appropriate flags to perform a *numeric* sort on the *key field* containing the count for each word in *descending* order. Pipe the output into head to show the top 3 words in King Lear and the number of times they appear

```
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:apache_beam.io.filebasedsink:Starting finalize write threads with num_shards: 1 (skipped: 0), batches: 1, num_threads: 1
INFO:apache_beam.io.filebasedsink:Renamed 1 shards in 0.01 seconds.
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$ cat outputs-*--of-* | wc -w
9568
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$ cat outputs-*--of-* | sort -rnk2 | head -n 3
the: 786
i: 622
and: 594
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$
```

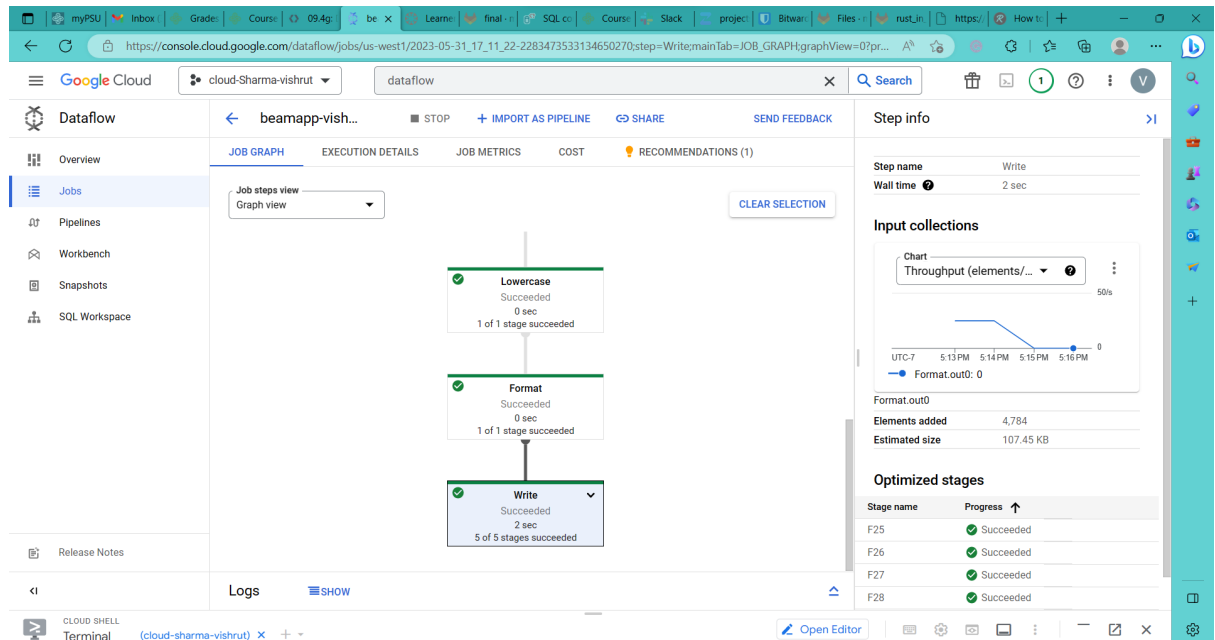
- Use the previous method to show the top 3 words in King Lear, case-insensitive, and the number of times they appear.

```
0x7f1b40c9f200> for environment ref Environment default_environment_1 (beam:env:embedded_python:vi, b'')
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:root:Default Python SDK image for environment is apache/beam_python3.9_sdk:2.48.0
INFO:apache_beam.io.filebasedsink:Starting finalize write threads with num_shards: 1 (skipped: 0), batches: 1, num_threads: 1
INFO:apache_beam.io.filebasedsink:Renamed 1 shards in 0.01 seconds.
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$ cat outputs-*--of-* | sort -rnk2 | head -n 3
the: 908
and: 738
i: 622
(env) vishrut@cloudshell:~/training-data-analyst/courses/machine_learning/deepdive/04_features/dataflow/python (cloud-sharma-vishrut)$
```

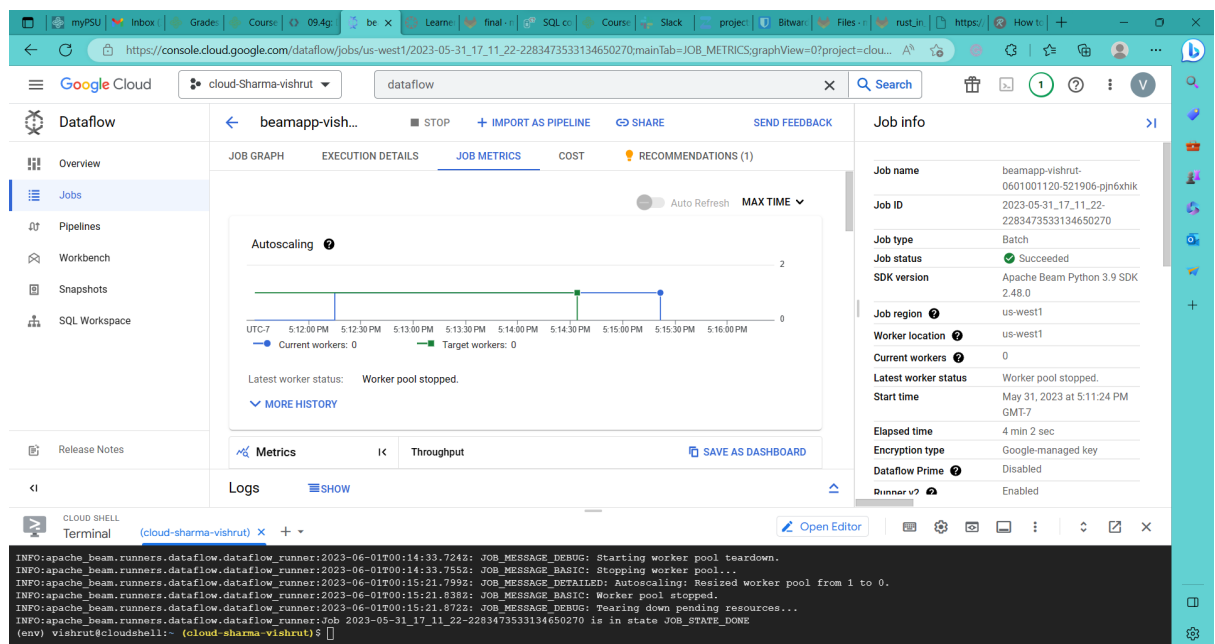
## 9. Run code using Dataflow runner

- The part of the job graph that has taken the longest time to complete.

**Answer:** Write took 2 seconds



- The autoscaling graph showing when the worker was created and stopped.



- Examine the output directory in Cloud Storage. How many files has the final write stage in the pipeline created?

**Answer:** 6 files

## 12. View raw data from PubSub

- Take a screenshot listing the different fields of this object.

```
(env) vishrut@cloudshell:~ (cloud-sharma-vishrut)$ gcloud pubsub subscriptions create taxisub \
--topic=projects/pubsub-public-data/topics/taxirides-realtime
Created subscription [projects/cloud-sharma-vishrut/subscriptions/taxisub].
(env) vishrut@cloudshell:~ (cloud-sharma-vishrut)$ gcloud pubsub subscriptions pull taxisub --auto-ack
DATA: ("ride_id":2544244e-b237-4755-bf8b-501d045e9d49","point_id":327,"latitude":40.75836,"longitude":-73.99269000000001,"timestamp":"2023-05-31T20:35:54.82612-04:00","meter_reading":
10.384876,"meter_increment":0.031758035,"ride_status":"enroute","passenger_count":1)
MESSAGE ID: 7867434769498327
ORDERING KEY:
ATTRIBUTES: ts=2023-05-31T20:35:54.82612-04:00
DELIVERY_ATTEMPT:
ACK_STATUS: SUCCESS
(env) vishrut@cloudshell:~ (cloud-sharma-vishrut)$
```

## 14. Run Dataflow job from template

- Take a screenshot of the pipeline that includes its stages and the number of elements per second being handled by individual stages.

## 15. Query data in BigQuery

- Take a screenshot showing the number of passengers and the amount paid for the first ride

Row	point_idx	latitude	longitude	timestamp	meter_reading	meter_increment	ride_status	passenger_count
1	aa8...	132	40.7411500...	2023-06-01 00:38:06.712140 U...	4.185992	0.03171206	enroute	1
2	c99...	1588	40.6441300...	2023-06-01 00:38:06.695010 U...	28.387762	0.017876424	enroute	1
3	e5e...	216	40.77743	2023-06-01 00:38:16.808850 U...	7.6712446	0.03551502	enroute	2
4	:19...	603	40.74166	2023-06-01 00:38:16.855830 U...	12.567659	0.02084189	enroute	2
5	:1d9...	226	40.77982	2023-06-01 00:38:17.215310 U...	6.9669175	0.030827068	enroute	2
6	6cd...	248	40.7108600...	2023-06-01 00:38:18.197350 U...	5.860274	0.023630137	enroute	3
7	:59...	434	40.7465200...	2023-06-01 00:38:18.648950 U...	8.731244	0.020118074	enroute	1
8	d78...	976	40.7211600...	2023-06-01 00:38:19.290740 U...	24.31			5
9	764...	892	40.7487000...	2023-06-01 00:38:18.930530 U...	21.75			1
10	:1f...	45	40.74629	2023-06-01 00:38:19.731600 U...				1
11	:546...	2584	40.79905	2023-06-01 00:38:18.915150 U...	55.3			1
12	106...	170	40.7703400...	2023-06-01 00:38:20.822510 U...	9.63			1
13	2e1...	2	40.76881	2023-06-01 00:38:14.943310 U...	0.0530			1
14	4e7...	458	40.78477	2023-06-01 00:38:20.174970 U...	11.70			1
15	a44...	61	40.74645	2023-06-01 00:38:12.201840 U...	3.013022			1

- Take a screenshot showing the estimated number of rows in the table.

### Streaming buffer statistics

Estimated size	287.05 MB
Estimated rows	1,997,234
Earliest entry time	May 31, 2023, 5:39:58 PM UTC-7

PERSONAL HISTORY

PROJECT HISTORY

- Take a screenshot showing the per-minute number of rides, passengers, and revenue for the data collected

Query results

Row	minute	total_rides	total_passengers	total_revenue
1	17:37	4	5	0.0
2	17:38	413	685	5854.430003400...
3	17:39	437	767	6224.6299922
4	17:40	467	814	6936.830019499...
5	17:41	459	743	6040.040003100...
6	17:42	476	769	6924.410018600...
7	17:43	443	771	6255.1100022
8	17:44	442	819	6708.740005400...
9	17:45	449	793	6433.729998500...
10	17:46	466	797	6517.219991500...
11	17:47	456	757	6714.039982199...
12	17:48	19	22	233.9000005

## 16. Data visualization

- Take a screenshot showing the plot for your data for your lab notebook

